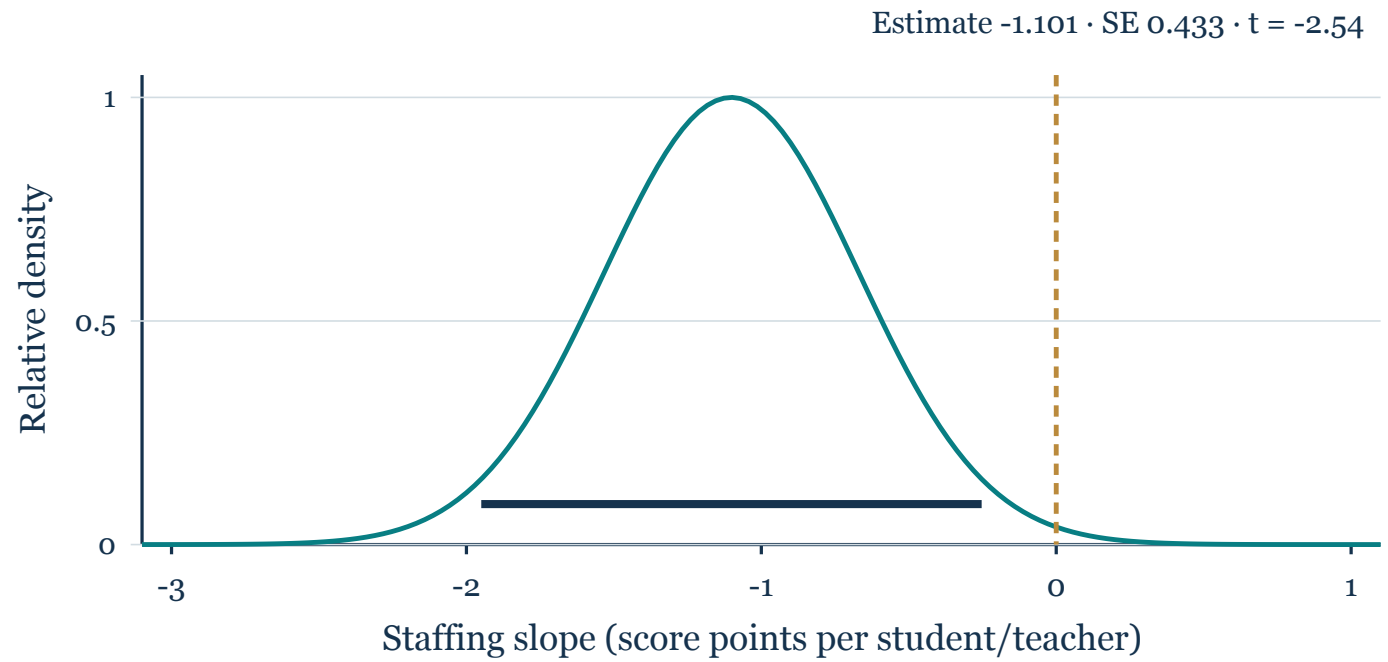
Is the staffing relationship large enough to matter?

A decision needs magnitude and uncertainty.



Actual staffing estimate and robust SE. The reference curve is a large-sample normal approximation, not the distribution of district outcomes.

Recall which coefficient the model estimates

Keep its units and adjustment fixed.

> The adjusted staffing slope is -1.101 points per additional student per teacher.

> Before calculating a test, state the exact target.

An outcome distribution is not a sampling distribution

The random object changes.

> A distribution describes the possible values of a random quantity and how often they occur.

> These are different objects.

A standard error estimates the estimator’s spread

It is not the standard deviation of the outcome.

> $$t = \frac{\widehat{\beta}_j - \beta_{j,0}}{\mathrm{SE}(\widehat{\beta}_j)}$$

> $\beta_{j,0}$: proposed null value; t: gap measured in SE units

Build a confidence interval step by step

Combine the estimate with a stated reference rule.

> $\widehat{\beta} \pm 1.96\,\mathrm{SE}(\widehat{\beta})$

Confidence interval

Apply the idea before reading the solution.

Practice

An estimate is 0.30 with robust SE 0.12. Use the large-sample 1.96 rule to form a 95% interval and test a zero effect.

Confidence interval: explanation

Connect the calculation to the empirical claim.

> $0.30 \pm 1.96 \times 0.12$ gives [0.0648,0.5352]. The t ratio is 2.5, so reject zero at approximately 5% under the relevant large-sample assumptions. This says nothing by itself about the economic importance or causal validity of the estimate.

State a null before calculating a p-value

The evidence is evaluated under a specific claim.

> A null hypothesis is a specified population statement, such as a zero staffing slope.

> A p-value is the probability, under the null and test assumptions, of a statistic at least as extreme as the observed statistic.

A rejection is not an effect-size statement

Statistical significance and practical value answer different questions.

> A statistically significant result passes a chosen testing threshold.

> For a staffing decision, translate the slope into a feasible ratio change and compare the implied learning difference with costs and relevant alternatives.

Normal-error exact inference and robust approximations differ

Know which argument supports the reference distribution.

> Under the classical linear model with normally distributed errors and the other required conditions, certain t and F procedures have exact finite-sample reference distributions.

> Robust inference generally relies on a large-sample approximation instead.

Unequal error spread is heteroskedasticity

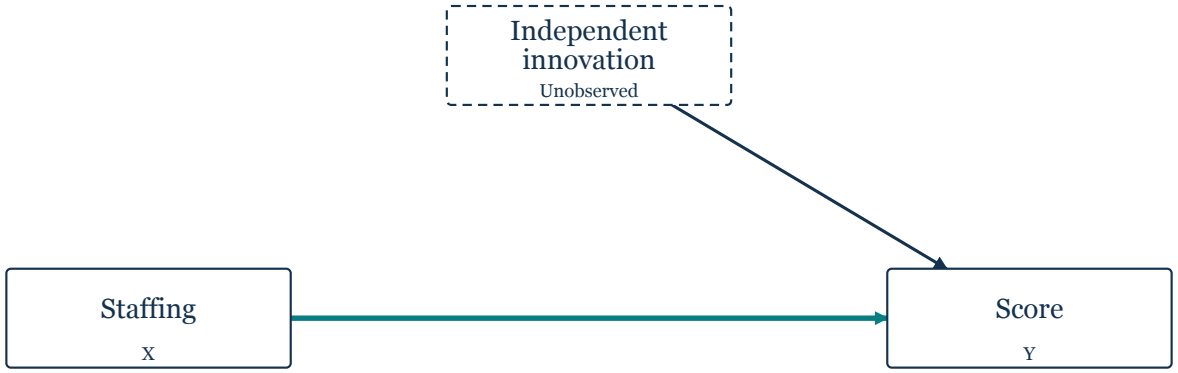
It concerns the error variance at different regressor values.

> $E(u_i \mid X_i) = 0$ concerns the mean

> $\mathrm{Var}(u_i \mid X_i)$ concerns spread

A DAG does not encode the size of the noise

Use a mean-zero innovation with a regressor-dependent scale.



1 open directed route · 0 open alternatives · gold: conditioned on

Assumed causal structure · dashed boxes: unobserved · gold boxes: conditioned on

$$Y = a + bX + \sigma(X)\varepsilon$$

$$E(u \mid X) = 0, \quad \mathrm{Var}(u \mid X) = \sigma(X)^2$$

Compare conventional and robust output

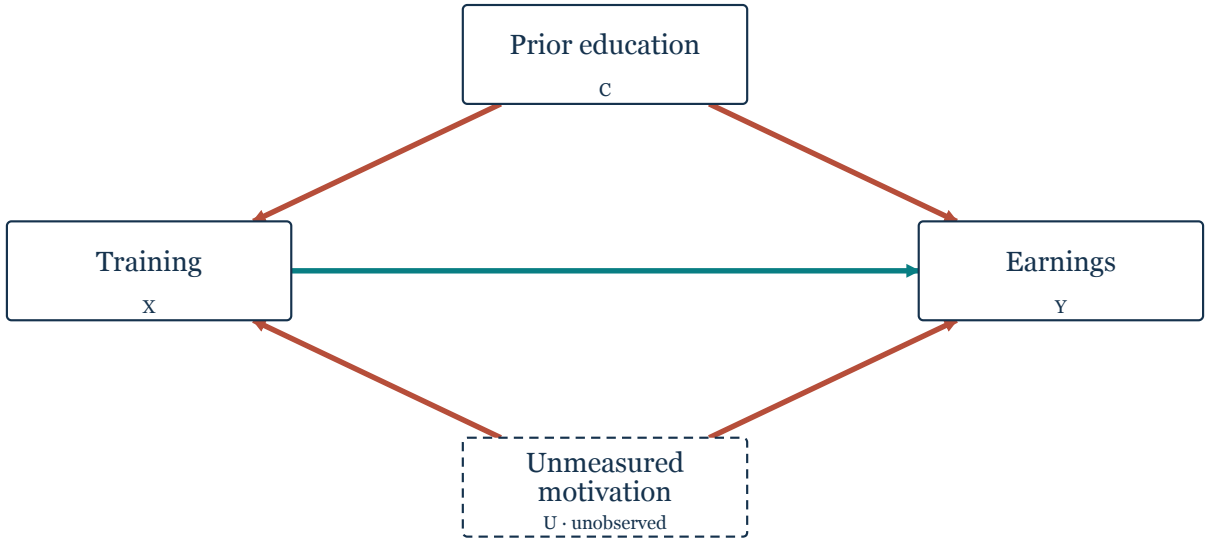
Same coefficients, different estimated uncertainty.

```
. regress score str el, vce(robust)

------------------------------------------------------------------------------
             |               Robust
       score | Coefficient  std. err.      t    P>|t|     [95% conf. interval]
-------------+----------------------------------------------------------------
         str |  -1.101296   .4328472    -2.54   0.011     -1.95213   -.2504616
          el |  -.6497768   .0310318   -20.94   0.000     -.710775   -.5887786
       _cons |   686.0322   8.728224    78.60   0.000     668.8754     703.189
------------------------------------------------------------------------------
```

Do not choose an SE by the resulting significance

Match it to the sampling and dependence structure.



1 open directed route · 2 open alternatives · gold: conditioned on

Assumed causal structure · dashed boxes: unobserved · gold boxes: conditioned on

Robust does what?

Apply the idea before reading the solution.

Practice

After changing regress y x to regress y x, vce(robust), which point estimate changes? Does this solve omitted-variable bias?

Robust does what?: explanation

Connect the calculation to the empirical claim.

> On the same sample the OLS coefficient does not change. The estimated covariance matrix changes, so SEs, intervals and tests can change. Heteroskedasticity-robust inference does not remove omitted-variable bias or automatically allow serial dependence.

Test a substantively meaningful value

A zero null is not the only relevant question.

```
. test str=-1

 ( 1)  str = -1

       F(  1,   417) =    0.05
            Prob > F =    0.8151
```

Compute uncertainty for the policy contrast

Scale the SE along with the estimate.

```
. lincom -2*str

------------------------------------------------------------------------------
       score | Coefficient  Std. err.      t    P>|t|     [95% conf. interval]
-------------+----------------------------------------------------------------
         (1) |   2.202592   .8656943     2.54   0.011     .5009232     3.90426
------------------------------------------------------------------------------
```

$$\mathrm{SE}(-2\widehat{\beta}_1) = 2\,\mathrm{SE}(\widehat{\beta}_1)$$

A confidence interval can contain an important effect

Failing to reject is not proof of no effect.

> If an interval includes zero and also includes effects large enough to matter, the evidence may be too imprecise for the decision.

> Report the point estimate and interval, explain which practical values remain compatible with the model and discuss what data would improve precision.

More observations help only under the relevant conditions

The information gain depends on what the observations add.

> Illustrative rate: SE $\propto n^{-1/2}$

Avoid repeated specification searches for a preferred p-value

The search changes what an isolated test can mean.

> If an analyst tries many outcomes, control sets and sample restrictions and reports only the most significant result, the reported p-value does not describe the whole search.

> Keep a clear account of the main question and intended specification.

Write an uncertainty statement for a policy reader

Include the comparison, units and conditions.

> A concise result sentence names the adjusted comparison, gives a feasible policy contrast and reports its uncertainty method.

> For example, describe a fitted two-student ratio difference at fixed language composition and give its robust interval in score points.

Economic versus statistical size

Apply the idea before reading the solution.

Practice

A precisely estimated policy effect is 0.001 outcome units. Is a small p-value enough to recommend the policy?

Economic versus statistical size: explanation

Connect the calculation to the empirical claim.

> No. Report the units, realistic policy contrast, interval, policy cost and relevant outcome scale. Statistical detectability alone neither establishes causal validity nor a worthwhile benefit.

Check the reference distribution and scale

An interval is meaningful only for a defined quantity.

> Before interpreting an interval, identify whether it concerns a coefficient, a fitted contrast, an average outcome or an individual prediction.

> Also check whether the displayed interval uses a normal approximation or a t reference and whether the covariance option matches the observations.

Use the evidence to identify the next useful step

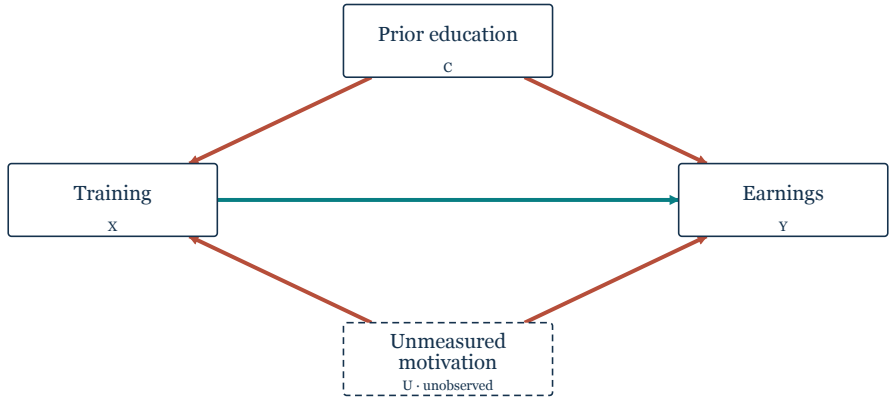
Precision, fit and design can require different improvements.

> An imprecise but credible comparison may benefit from more informative observations.

> These are different diagnoses.

DAG practice: uncertainty is not adjustment

Trace the routes before choosing a control.



1 open directed route · 2 open alternatives · gold: conditioned on

Assumed causal structure · dashed boxes: unobserved · gold boxes: conditioned on

Practice

Select prior education as a control in the training diagram. Which alternative route remains? If the analyst changes conventional SEs to heteroskedasticity-robust SEs, does that close the route or change the training coefficient on an identical specification and sample?

DAG practice: uncertainty is not adjustment: explanation

Explain what conditioning changes in this stated graph.

> The route through unmeasured motivation remains open after conditioning on education. The common-cause arrows are properties of the proposed causal story, not of the covariance formula. Switching the SE option does not block them. On an unchanged model and sample, the OLS coefficient remains the same while the estimated covariance, intervals and tests can change. Robust inference addresses uncertainty under appropriate conditions. A credible causal comparison needs an additional argument about motivation or a design that addresses its role.

Report magnitude and uncertainty together

Reproduce the evidence, then explain its economic meaning.

Practice

Run meeting_04.do. Use the robust staffing coefficient and SE to construct the approximate 95% interval with the 1.96 rule. Then calculate the estimate and SE for a two-student reduction. Explain how a decision about practical benefits differs from a test of a zero coefficient.

Report magnitude and uncertainty together: explanation

A reproducible result needs both a calculation and an interpretation.

> The robust staffing estimate is -1.101296, with SE 0.432847. The approximate interval is [-1.9497, -0.2529] score points per additional student per teacher. Multiplying the slope by minus two gives 2.2026 points; its SE is 0.8657, because scaling by a constant scales the SE by its absolute value. These calculations quantify uncertainty under the model and inference assumptions. Stata's finite-degree-of-freedom interval can differ slightly from the 1.96 approximation. A recommendation also needs a credible causal comparison, the outcome's educational importance, staffing costs and the relevant population. Rejecting a zero slope does not supply those additional pieces of the decision.

Next, let the model represent percentages and curvature

The interpretation depends on how variables enter.

> We have connected a coefficient with its uncertainty and separated detectability from practical size.

> The next lesson changes the equation itself.